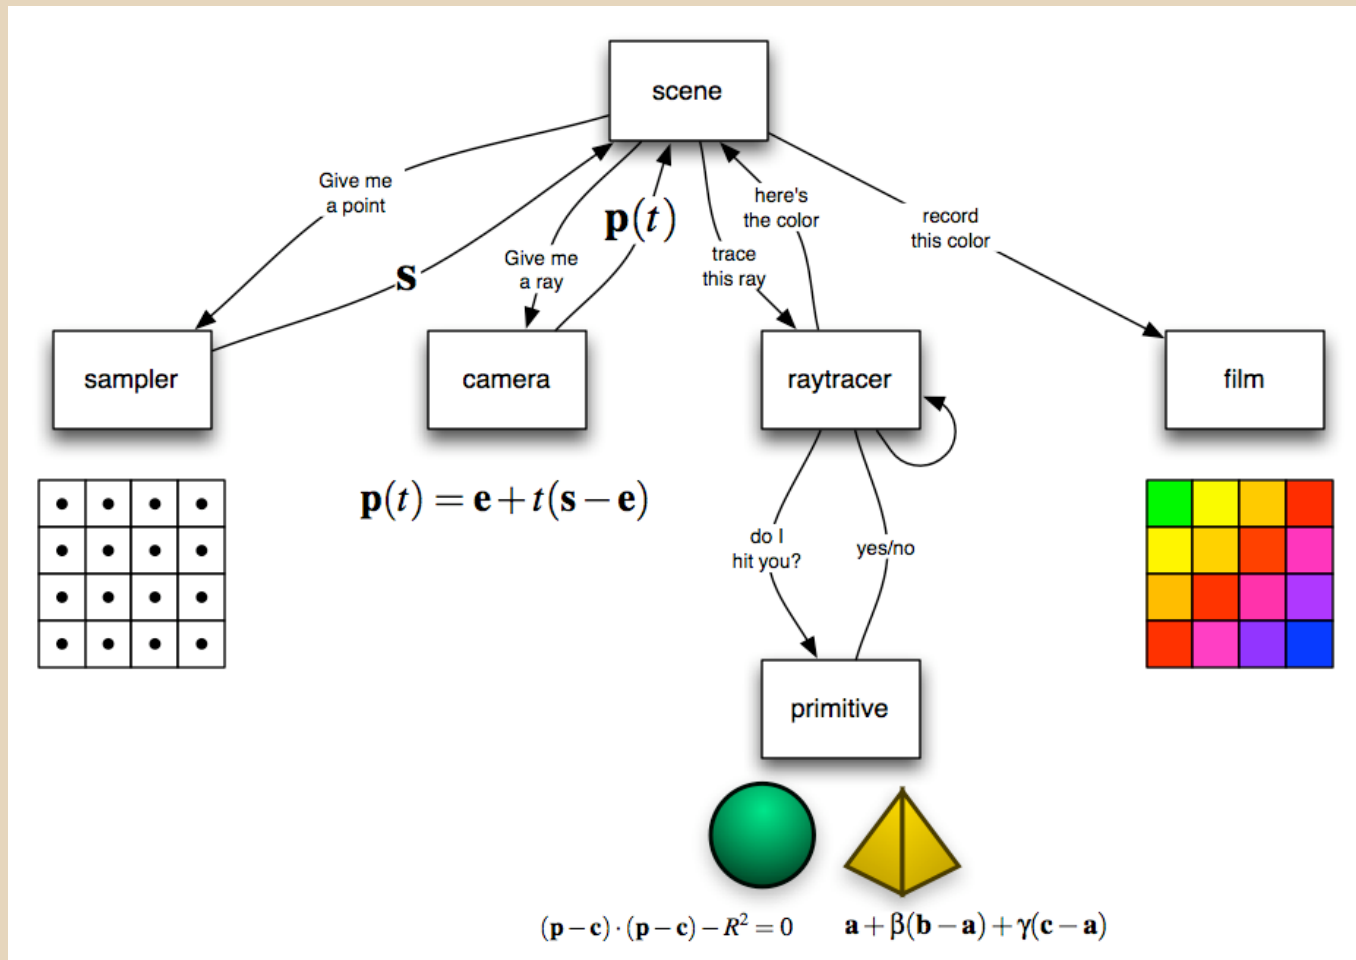


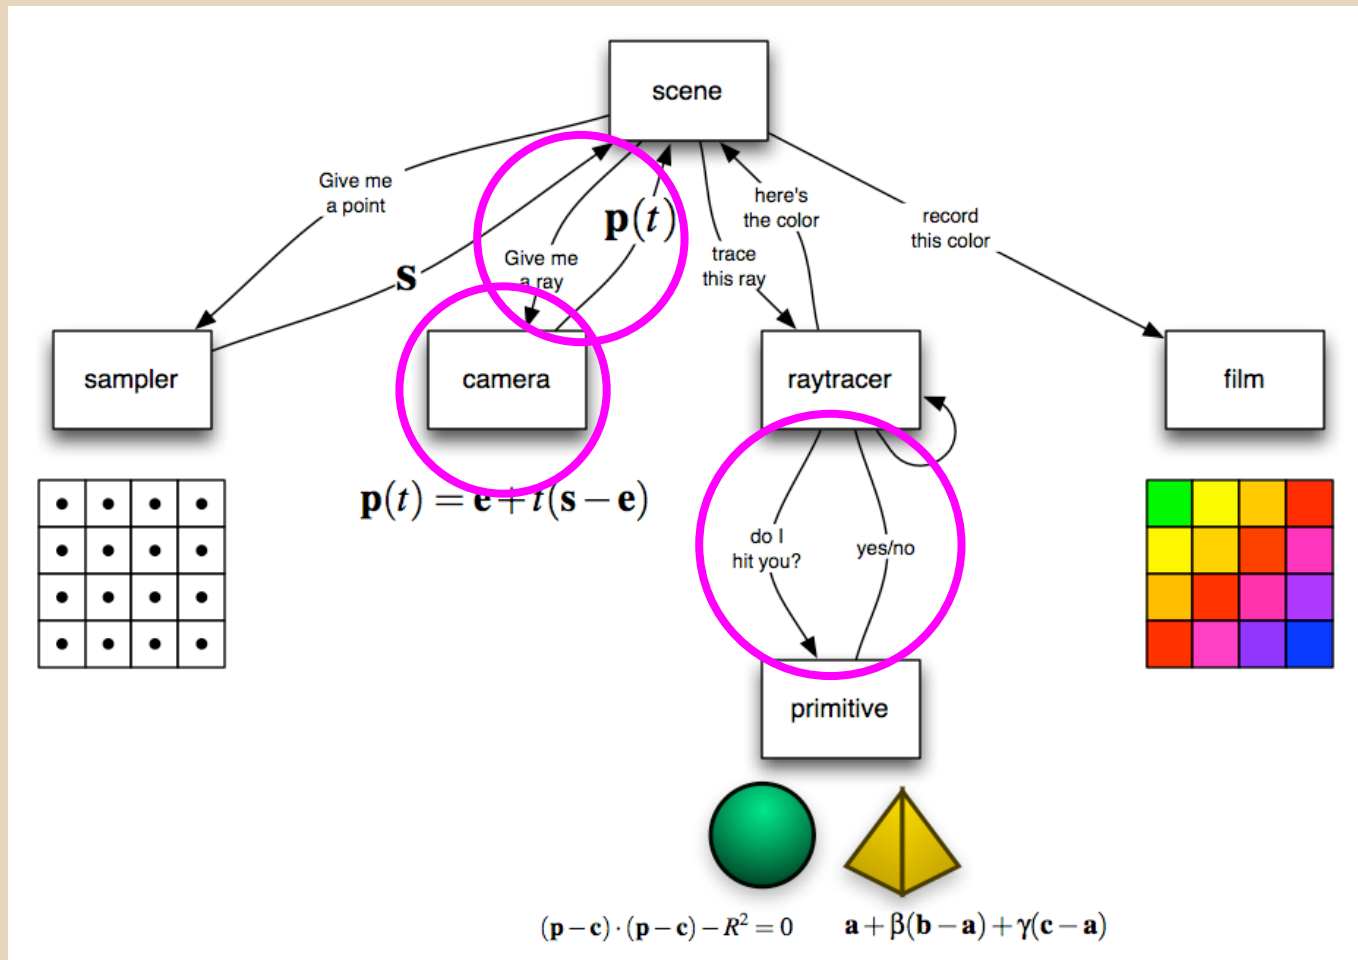
Raytracer Topics: Cameras and Raycasting

Adam Mally
CIS 560 Spring 2015
University of Pennsylvania

Basic Structure of a Raytracer




Today's Topics



Viewing a scene

- Concatenate three matrices with your geometry to view your scene
- $\text{Projection_Mat} * \text{View_Mat} * \text{Model_Mat} * \text{Geometry}$
- Model matrix:
 - Object space \rightarrow world space
 - Scale, rotation, and translation in any combination
- View matrix:
 - Set up your camera's position and orientation
- Projection matrix:
 - Project objects from 3D space into the screen's 2D space
 - Screen space is in Normalized Device Coordinates



No matrices


Model matrix

View matrix

Projection matrix

Camera: View matrix parameters

- Position in world coordinate system (commonly referred to as the **eye position**)
- (Optional) A point at which to look, commonly referred to as the **reference point**
- **Look direction (F)** - which way is the camera facing? Represents the local Z-axis
- An “**up direction**” in world coordinates, used for computing the camera’s local axes
- **Local right (R)** - A direction in world coordinates that represents the direction that is “rightward” in the camera’s local coordinates. Represents local X-axis
 - Perpendicular to the **look direction**
- **Local up (U)** - A direction in world coordinates that represents the direction that is “upward” in the camera’s local coordinates. Represents local Y-axis
 - Perpendicular to both **local right** and **look direction**




Camera: Axis computation

$F = \text{normalize}(\text{ref} - \text{eye})$

$R = \text{normalize}(F \times \text{world_up})$

$U = \text{normalize}(R \times F)$



Camera: View matrix computation


- View matrix transforms geometry from world space into camera space
 - The camera itself is not moving; the scene's geometry is transformed in the opposite direction and orientation of how you want the camera to move
 - Strictly speaking, the “camera” does not exist; you are always transforming the geometry in your scene
- View matrix is composed of two sub-matrices
 - Orientation matrix (O)
 - Translation matrix (T)

$$\bullet \quad O = \begin{bmatrix} R_x & R_y & R_z & 0 \\ U_x & U_y & U_z & 0 \\ F_x & F_y & F_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad \bullet \quad T = \begin{bmatrix} 1 & 0 & 0 & -eye_x \\ 0 & 1 & 0 & -eye_y \\ 0 & 0 & 1 & -eye_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- View matrix = $O * T$

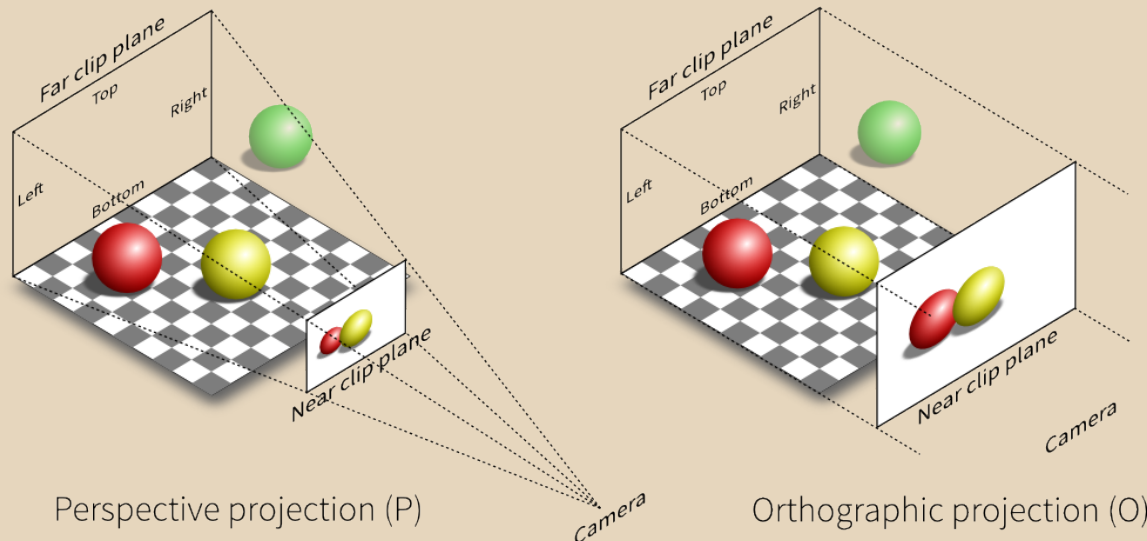
Camera: Projection matrix

- Projection matrix transforms geometry from camera space into screen space (remember: NDC)
- Built from several camera components
 - Field of view - the angle that determines how quickly the frustum grows as it extends from the camera
 - Aspect ratio - screen width / screen height
 - Near clip and far clip planes
 - Projection type - perspective or orthographic



Camera: Viewing frustums

- Geometry outside the frustum is not visible to the camera
- Frustum shape affects how geometry is projected
- Since perspective projection frustums grow in size as they extend from the camera, geometry that is further from the camera appears to be smaller than geometry closer to the camera
 - The geometry gets further from the frustum bounds as they grow
- Orthographic frustums do not grow with distance, so a sphere 1 foot away and a sphere 10 feet away will look the same



Camera: Orthographic projection

- Given a near clip, far clip, top clip, bottom clip, and aspect ratio:

$$\begin{aligned} \text{bottom} &= -\text{top} \\ \text{right} &= \text{top} \times \text{aspect} \\ \text{left} &= -\text{right} \end{aligned}$$
$$P = \begin{pmatrix} \frac{2}{\text{right} - \text{left}} & 0 & 0 & -\frac{\text{right} + \text{left}}{\text{right} - \text{left}} \\ 0 & \frac{2}{\text{top} - \text{bottom}} & 0 & -\frac{\text{top} + \text{bottom}}{\text{top} - \text{bottom}} \\ 0 & 0 & \frac{1}{\text{far} - \text{near}} & -\frac{\text{near}}{\text{far} - \text{near}} \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

- Maps (left, bottom, near) to (-1, -1, 0) and (right, top, far) to (1,1,1)

Camera: Orthographic projection

- Let's do some matrix multiplication to see how the clipping planes work
- Near clip: 1
- Far clip: 10
- Top: 10
- Aspect: 1:1
- Camera position: [0,0,0]
- Look vector: [0,0,1]

$$\begin{aligned}
 \text{bottom} &= -\text{top} \\
 \text{right} &= \text{top} \times \text{aspect} \\
 \text{left} &= -\text{right}
 \end{aligned}$$

$$P = \begin{pmatrix} \frac{2}{\text{right} - \text{left}} & 0 & 0 & -\frac{\text{right} + \text{left}}{\text{right} - \text{left}} \\ 0 & \frac{2}{\text{top} - \text{bottom}} & 0 & -\frac{\text{top} + \text{bottom}}{\text{top} - \text{bottom}} \\ 0 & 0 & \frac{1}{\text{far} - \text{near}} & -\frac{\text{near}}{\text{far} - \text{near}} \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

$$P = \begin{pmatrix} 0.1 & 0 & 0 & 0 \\ 0 & 0.1 & 0 & 0 \\ 0 & 0 & 0.111 & -0.111 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

- Note how the upper-right element and the element below that always zero out, since they're $(\text{right} + \text{left})/(\text{numbers})$ and $(\text{top} + \text{bottom})/(\text{numbers})$, which are equivalent to $(\text{right} - \text{right})/(\text{numbers})$ and $(\text{top} - \text{top})/(\text{numbers})$

Camera: Orthographic projection

$$\begin{vmatrix} 0.1 & 0 & 0 & 0 \\ 0 & 0.1 & 0 & 0 \\ 0 & 0 & 0.111 & -0.111 \\ 0 & 0 & 0 & 1 \end{vmatrix} * \begin{vmatrix} -2 \\ 5 \\ 0.5 \\ 1 \end{vmatrix} = \begin{vmatrix} 0.1*-2 + 0*5 + 0*0.5 + 0*1 \\ 0*-2 + 0.1*5 + 0*0.5 + 0*1 \\ 0*-2 + 0*5 + 0.111*0.5 + -0.111*1 \\ 0*-2 + 0*5 + 0*0.5 + 1*1 \end{vmatrix}$$

$$= \begin{vmatrix} -0.2 \\ 0.5 \\ -0.0555 \\ 1 \end{vmatrix}$$

- Since the Z range of the screen is [0,1], we cannot see this point as its Z coordinate is negative
- Remember, though: the X and Y ranges are [-1, 1]
- Similarly, a point that is past the far clip plane would have a Z coordinate greater than 1 after being transformed by this matrix
- Overall, any point in space outside the frustum is mapped to a point outside the range <-1,-1,0> to <1,1,1>

Camera: Perspective projection

- Given a near clip, far clip, FOV, and aspect ratio:

$$\begin{aligned} \text{top} &= \text{near} \times \tan(\text{FOV}/2) \\ \text{bottom} &= -\text{top} \\ \text{right} &= \text{top} \times \text{aspect} \\ \text{left} &= -\text{right} \end{aligned}$$
$$P = \begin{pmatrix} \frac{2 \times \text{near}}{\text{right} - \text{left}} & 0 & -\frac{\text{right} + \text{left}}{\text{right} - \text{left}} & 0 \\ 0 & \frac{2 \times \text{near}}{\text{top} - \text{bottom}} & -\frac{\text{top} + \text{bottom}}{\text{top} - \text{bottom}} & 0 \\ 0 & 0 & \frac{\text{far}}{\text{far} - \text{near}} & -\frac{\text{far} \times \text{near}}{\text{far} - \text{near}} \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

- Maps (left, bottom, near) to (-1, -1, 0) and (right, top, near) to (1,1,0)

Camera: Perspective projection

- Let's do some matrix multiplication to see how this scales geometry
- Near clip: 1
- Far clip: 10
- FOV: 90
- Aspect: 1:1
- Camera position: [0,0,0]
- Look vector: [0,0,1]

$$\begin{aligned}\text{top} &= \text{near} \times \tan(\text{FOV}/2) \\ \text{bottom} &= -\text{top} \\ \text{right} &= \text{top} \times \text{aspect} \\ \text{left} &= -\text{right}\end{aligned}$$

$$P = \begin{pmatrix} \frac{2 \times \text{near}}{\text{right} - \text{left}} & 0 & -\frac{\text{right} + \text{left}}{\text{right} - \text{left}} & 0 \\ 0 & \frac{2 \times \text{near}}{\text{top} - \text{bottom}} & -\frac{\text{top} + \text{bottom}}{\text{top} - \text{bottom}} & 0 \\ 0 & 0 & \frac{\text{far}}{\text{far} - \text{near}} & -\frac{\text{far} \times \text{near}}{\text{far} - \text{near}} \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

$$P = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1.111 & -1.111 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

Camera: Perspective projection

$$\begin{vmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1.111 & -1.111 \\ 0 & 0 & 1 & 0 \end{vmatrix} * \begin{vmatrix} -1 \\ 1.5 \\ 2 \\ 1 \end{vmatrix} = \begin{vmatrix} 1*-1 + 0*1.5 + 0*2 + 0*1 \\ 0*-1 + 1*1.5 + 0*2 + 0*1 \\ 0*-1 + 0*1.5 + 1.111*2 + -1.111*1 \\ 0*-1 + 0*1.5 + 1*2 + 0*1 \end{vmatrix}$$

$$= \begin{vmatrix} -1 \\ 1.5 \\ 1.111 \\ 2 \end{vmatrix} = \begin{vmatrix} -0.5 \\ 0.75 \\ 0.556 \\ 1 \end{vmatrix}$$

- In order to maintain homogeneous coordinates, we must divide the entire vector by its w component so w remains 1
- Note how the last row of the matrix has a 1 in the Z component, which causes the w component of the resultant vector to scale with the Z coordinate of the point we're transforming

Camera matrices

What would we see if we were to draw a 1x1x1 cube centered at the origin with neither view nor projection matrices?


1




2



3




4




What is raycasting?

- Creating a line that passes through the viewing frustum and travels from the eye to some endpoint on a slice of the frustum (e.g. the far clip plane)
- The line's endpoint is determined by the pixel on our screen from which we want to raycast



Normalized device coordinates

- Recall that your GL window ranges from -1 to 1 on both the X and Y axes
- We can convert to NDC from any given pixel




- $sx = (2 * px / scr_width) - 1$
- $sy = 1 - (2 * py / scr_height)$
- px and py are the given pixel's x and y coordinates

Recap: Camera axes

$\mathbf{F} = \text{normalize}(\text{ref} - \text{eye})$


$\mathbf{R} = \text{normalize}(\mathbf{F} \times \text{world_up})$

$\mathbf{U} = \text{normalize}(\mathbf{R} \times \mathbf{F})$




Screen point to world point

- $len = |ref - eye|$
- $V = U * len * \tan(\alpha)$
- $H = R * len * \text{aspect} * \tan(\alpha)$
- $\alpha = FOVY / 2$
- $p = ref + sx * H + sy * V$
- sx, sy are in NDC




Getting a ray from the world point

- $\text{ray_origin} = \text{eye}$
- $\text{ray_direction} = \text{normalize}(p - \text{eye})$
- $\text{Arbitrary point on ray} = \text{eye} + t * \text{ray_direction}$




Ray-polygon intersection

- Most common intersection test is ray-triangle
- We'll also cover ray-sphere and ray-cube
 - All three are commonly used in basic raytracer testing




Frames of reference

- Before you try to test a ray against primitive geometry, you must first transform the ray so from its perspective, the geometry in question is primitive
- Simply transform the ray's direction and origin by the *inverse* of the geometry's model matrix




Frames of reference

- Similarly, make sure you transform the results of your intersection test back into world space (e.g. the point of intersection, the surface normal at the intersection, etc.)



Ray-sphere intersection

- Sphere defined as $(x_s - x_c)^2 + (y_s - y_c)^2 + (z_s - z_c)^2 = r_s^2$
 - Sphere center = $\langle x_c \ y_c \ z_c \rangle$
 - All points on the sphere surface = $\langle x_s \ y_s \ z_s \rangle$
 - r_s is the sphere's radius
- Ray defined as: $R_\theta + t * R_d$
 - $R_\theta = \langle x_\theta \ y_\theta \ z_\theta \rangle$
 - $R_d = \langle x_d \ y_d \ z_d \rangle$
 - t is a parameterization of R_d (i.e. a float)



Ray-sphere intersection

- Substitute $\langle x_s y_s z_s \rangle$ for the ray equation:

$$(x_\theta + t*x_d - x_c)^2 + (y_\theta + t*y_d - y_c)^2 + (z_\theta + t*z_d - z_c)^2 = r_s^2$$

- Can also be written as:

- $At^2 + Bt + C = 0$

- $A = x_d^2 + y_d^2 + z_d^2$

- $B = 2(x_d(x_\theta - x_c) + y_d(y_\theta - y_c) + z_d(z_\theta - z_c))$

- $C = (x_\theta - x_c)^2 + (y_\theta - y_c)^2 + (z_\theta - z_c)^2 - r_s^2$

- Note that we now have a quadratic equation


- We can solve for t using the quadratic formula!

Ray-sphere intersection

- $t_0, t_1 = (-B \pm \sqrt{B^2 - 4AC}) / (2A)$
 - t_0 is for the - case and t_1 is for the + case
- Remember: if the discriminant is negative, then there is no real root and therefore no intersection
 - Discriminant = $B^2 - 4AC$
- If t_0 is positive, then we're done. If not, then compute t_1 .

Ray-sphere intersection

- Once we have t , we can plug it into our ray equation to find the closest point of intersection on our sphere.
 - If all we care about is whether or not we hit the sphere, we can just check:
 - $\text{near_clip} < t < \text{far_clip}$
- $P = R_0 + t * R_d$




Converting from local to world space

- It is important to remember that the t value we have found is in **object space** (i.e. only valid for an untransformed sphere)
- To find a t value in world space, we can use our object-space t to find the point of intersection P_o on our unit sphere
- We can then transform P_o by our sphere's transformation matrix to find the point of intersection in world space P_w
- We can use P_w to compute our world-space t value as follows:
 - `world_t = length(Pw - camera.eye)`
 - Remember that `camera.eye` is the origin of your untransformed ray in world space


Ray-cube intersection

- Begin by storing t_{near} = -infinity and t_{far} = infinity
- For each pair of planes associated with the X, Y, and Z axes (the example uses the X “slab”):
 - If x_d is 0, then the ray is parallel to the X slab, so
 - If $x_0 < x_{\min}$ or $x_0 > x_{\max}$ then we miss
 - $t_0 = (x_{\min} - x_0)/x_d$
 - $t_1 = (x_{\max} - x_0)/x_d$
 - If $t_0 > t_1$ then swap them
 - If $t_0 > t_{\text{near}}$ then $t_{\text{near}} = t_0$
 - If $t_1 < t_{\text{far}}$ then $t_{\text{far}} = t_1$
 - Repeat for Y and Z
 - If $t_{\text{near}} > t_{\text{far}}$ then we miss the box



Ray-plane intersection

- Plane defined as: $\text{dot}(\mathbf{N}, (\mathbf{P} - \mathbf{S})) = 0$
 - \mathbf{N} is the plane's normal
 - \mathbf{S} is some point on the plane
 - \mathbf{P} is the point of intersection
- Ray defined as: $\mathbf{R}_0 + t * \mathbf{R}_d$
- Substitute P for ray:
 - $\text{dot}(\mathbf{N}, (\mathbf{R}_0 + t * \mathbf{R}_d - \mathbf{S})) = 0$
- Solve for t:
 - $t = \text{dot}(\mathbf{N}, (\mathbf{S} - \mathbf{R}_0)) / \text{dot}(\mathbf{N}, \mathbf{R}_d)$



Point-in-triangle

- Use **barycentric coordinates** to test if P is within the bounds of a triangle
 - The **barycenter** of a triangle is its center of mass, often given unequal weighting to its vertices
- $S = \text{area}(P_1, P_2, P_3)$
- $S_1 = \text{area}(P, P_2, P_3)/S$
- $S_2 = \text{area}(P, P_3, P_1)/S$
- $S_3 = \text{area}(P, P_1, P_2)/S$
- Therefore, $P = S_1P_1 + S_2P_2 + S_3P_3$
- So, P is within the triangle if:
 - $0 \leq S_1 \leq 1$
 - $0 \leq S_2 \leq 1$
 - $0 \leq S_3 \leq 1$
 - $S_1 + S_2 + S_3 = 1$

